

Final-Notebook

May 18, 2020

1 Trends in College Basketball

1.0.1 By Evan Devore and Mat Steininger

Sports and data science have taken an explosive growth in the past few decades with prominent examples like Moneyball, SportVU in the NBA, and almost every major professional sport's teams hiring data analysts to gain a competitive edge. This is no surprise; teams will do anything they can to win more championships and increase their revenue. The sports analytics industry is expected to balloon to \$4.5 billion dollars in the next few years, according to <https://www.businesswire.com/news/home/20181205005823/en/Global-4.5-Billion-Sports-Analytics-Market-Forecasts>.

One such sport where data analytics has had a significant impact is professional basketball. However, with this tutorial, we are going to investigate the world of NCAA Men's College Basketball. Like professional sports, college teams use all sorts of analytics like game statistics all the way to movement data in order to gain a competitive edge. For more information on how analytics is affecting college basketball, see this article by Bleacher Report: <https://bleacherreport.com/articles/2807432-the-analytics-uprising-is-upon-college-basketball-how-it-could-alter-the-sport>.

In our tutorial, we are going to compare Final AP #1 teams across years of NCAA Men's College Basketball. For those unaware, the AP Polls are the NCAA standard for ranking teams and are decided by the media. At the end of each year, the final rankings show which teams are regarded as the best teams going into the NCAA tournament, and which squad is widely considered the best team of that regular season.

This tutorial will be organized based on the Data Science Pipeline that we learned in CMSC 320: 1. Data Collection 2. Data Processing 3. Exploratory Data Analysis and Visualization 4. Hypothesis Testing and Machine Learning 5. Final Thoughts

2 Part 1 - Data Collection

```
[2]: import requests
from bs4 import BeautifulSoup
import re
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
```

```

[3]: base_url = 'https://www.sports-reference.com'

# Data Frame for all #1 teams
top_teams = pd.DataFrame()

# Looping through every year from 1949 (Start of AP polling) to most recent
↳ final AP Poll in 2020
for year in range(1949, 2021):
    # Sports-Reference uses a simple year as the page format
    season_url = base_url + '/cbb/seasons/' + str(year) + '.html'
    season_request = requests.get(season_url).text
    soup = BeautifulSoup(season_request, features='lxml')

    # Find on the page where AP Final #1 is found and extract team page and name
    ap_final = soup.find(text=re.compile('AP Final #1')).parent.parent.parent.
    ↳ find("a")
    team_page_url = ap_final.get("href")
    team_name = ap_final.text

    # Now go into each team page and extra data
    team_page_request = requests.get(base_url + team_page_url).text
    soup = BeautifulSoup(team_page_request, features='lxml')
    table_team_data = soup.find(id="team_stats")
    if table_team_data is None:
        # No team data exists for this team for this season
        continue
    table_team_data = table_team_data.findAll("tr")

    # Create data frame to append to, year will be index
    year = team_page_url[team_page_url.rfind("/") + 1:team_page_url.find(".
    ↳ html")]
    team_data = pd.DataFrame(index=[year])
    team_data['team'] = team_name
    # Gets year from url

    # First row is labels for stats
    # Second row is teams stats for
    for entries in table_team_data[1].findAll("td"):
        stat = entries.get("data-stat")
        team_data[stat] = entries.text

    # Third row is teams stats against. Some years have different third row, so
    ↳ this takes that into account
    against = 2
    if len(table_team_data) > 3:
        against = -2
    for entries in table_team_data[against].findAll("td"):

```

```

stat = entries.get("data-stat") # + '_against'
team_data[stat] = entries.text

# Append our row to the ongoing list of AP #1 teams
top_teams = top_teams.append(team_data)

# Move team name to the front of the df
names = top_teams['team']
top_teams.drop(['team'], axis = 1, inplace = True)
top_teams.insert(0, 'team_name', names)

```

```
[4]: top_teams.columns
```

```
[4]: Index(['team_name', 'g', 'mp', 'fg', 'fga', 'fg_pct', 'ft', 'fta', 'ft_pct',
'orb', 'drb', 'trb', 'ast', 'stl', 'blk', 'tov', 'pf', 'pts',
'pts_per_g', 'opp_fg', 'opp_fga', 'opp_fg_pct', 'opp_ft', 'opp_fta',
'opp_ft_pct', 'opp_orb', 'opp_drb', 'opp_trb', 'opp_ast', 'opp_stl',
'opp_blk', 'opp_tov', 'opp_pf', 'opp_pts', 'opp_pts_per_g', 'fg2',
'fg2a', 'fg2_pct', 'fg3', 'fg3a', 'fg3_pct', 'opp_fg2', 'opp_fg2a',
'opp_fg2_pct', 'opp_fg3', 'opp_fg3a', 'opp_fg3_pct'],
dtype='object')
```

```
[5]: top_teams.tail()
```

```
[5]:
```

	team_name	g	mp	fg	fga	fg_pct	ft	fta	ft_pct	orb	...	\
2016	Kansas	38	7700	1092	2207	.495	601	843	.713	402	...	
2017	Villanova	36	7200	965	1948	.495	538	681	.790	316	...	
2018	Virginia	34	6825	848	1844	.460	340	451	.754	282	...	
2019	Duke	38	7625	1157	2418	.478	551	803	.686	495	...	
2020	Kansas	31	6225	851	1758	.484	411	616	.667	333	...	

	fg2_pct	fg3	fg3a	fg3_pct	opp_fg2	opp_fg2a	opp_fg2_pct	opp_fg3	opp_fg3a	\
2016	.533	304	728	.418	640	1474	.434	233	726	
2017	.592	311	843	.369	595	1211	.491	251	807	
2018	.501	247	645	.383	432	1009	.428	215	694	
2019	.580	278	903	.308	707	1571	.450	253	844	
2020	.553	199	578	.344	455	1065	.427	229	750	

	opp_fg3_pct
2016	.321
2017	.311
2018	.310
2019	.300
2020	.305

```
[5 rows x 47 columns]
```

2.0.1 Data Collection Tutorial

To start, we head over to the vast database of sports statistics over at <https://www.sports-reference.com>, which has both AP Final #1 teams and also that team's complete season statistics. The reason we chose AP Final #1 and not NCAA Tournament Champion is because the tournament is unpredictable, and some teams start playing well just at the right time. However, with AP Final #1 teams, they are a better representation of the best college basketball team over the course of the entire regular season, and thus have the statistics that represent the best team from that year.

Navigating and scraping data from this website is difficult because there is no single table with all the statistics we need. Firstly, each college basketball season has its own page which ends in year.html, so we can simply use the HTTP requests library in conjunction with BeautifulSoup to get the html of every season starting from when the first AP polling occurred in 1949. Then, by building the URL of each season, we can then find the AP Final #1 team and go to their respective page where all that season's data is nicely represented in a table. Certain years have less or no data because some statistics were not tracked until later seasons.

At this point, we have found an abundance of CBB data from each season's AP Poll 'Best Team' for each season starting at 1949, and loaded into the `top_teams` DataFrame. All of the columns for statistics from the team's opponent start with the `opp_` prefix. For example, the Total Rebounds statistic is kept under the `trb` column, but opponent total rebounds are kept under the `opp_trb` column. Additionally, the index of the row is the year in which the season was played. Furthermore, the final row in the DataFrame is the from 2019-2020 season.

3 Part 2 - Data Processing

Let's convert all of the columns to their proper datatypes. A float is necessary for columns that are percentages or per game values. Integers will suffice for the rest.

```
[6]: # Convert non-name columns to float or int
import numpy as np
from plotnine import *

top_teams.replace(r'^\s*$', np.nan, regex=True, inplace = True)
for column in top_teams:
    if not column == 'team_name':
        top_teams[column] = top_teams[column].fillna(-1)
        if ('pct' in column) or ('per_g' in column):
            top_teams[column] = top_teams[column].astype(float)
        else:
            top_teams[column] = top_teams[column].astype(int)
    print(column + " type is " + str(type(top_teams[column][-1])))
```

```
team_name type is <class 'str'>
g type is <class 'numpy.int64'>
mp type is <class 'numpy.int64'>
fg type is <class 'numpy.int64'>
fga type is <class 'numpy.int64'>
```

```

fg_pct type is <class 'numpy.float64'>
ft type is <class 'numpy.int64'>
fta type is <class 'numpy.int64'>
ft_pct type is <class 'numpy.float64'>
orb type is <class 'numpy.int64'>
drb type is <class 'numpy.int64'>
trb type is <class 'numpy.int64'>
ast type is <class 'numpy.int64'>
stl type is <class 'numpy.int64'>
blk type is <class 'numpy.int64'>
tov type is <class 'numpy.int64'>
pf type is <class 'numpy.int64'>
pts type is <class 'numpy.int64'>
pts_per_g type is <class 'numpy.float64'>
opp_fg type is <class 'numpy.int64'>
opp_fga type is <class 'numpy.int64'>
opp_fg_pct type is <class 'numpy.float64'>
opp_ft type is <class 'numpy.int64'>
opp_fta type is <class 'numpy.int64'>
opp_ft_pct type is <class 'numpy.float64'>
opp_orb type is <class 'numpy.int64'>
opp_drb type is <class 'numpy.int64'>
opp_trb type is <class 'numpy.int64'>
opp_ast type is <class 'numpy.int64'>
opp_stl type is <class 'numpy.int64'>
opp_blk type is <class 'numpy.int64'>
opp_tov type is <class 'numpy.int64'>
opp_pf type is <class 'numpy.int64'>
opp_pts type is <class 'numpy.int64'>
opp_pts_per_g type is <class 'numpy.float64'>
fg2 type is <class 'numpy.int64'>
fg2a type is <class 'numpy.int64'>
fg2_pct type is <class 'numpy.float64'>
fg3 type is <class 'numpy.int64'>
fg3a type is <class 'numpy.int64'>
fg3_pct type is <class 'numpy.float64'>
opp_fg2 type is <class 'numpy.int64'>
opp_fg2a type is <class 'numpy.int64'>
opp_fg2_pct type is <class 'numpy.float64'>
opp_fg3 type is <class 'numpy.int64'>
opp_fg3a type is <class 'numpy.int64'>
opp_fg3_pct type is <class 'numpy.float64'>

```

```
[7]: top_teams.head()
```

```

[7]:      team_name  g  mp   fg   fga fg_pct   ft   fta  ft_pct  orb  ...  \
1949  Kentucky  34  -1  903  2756  0.328  514  728   0.706  -1  ...

```

1951	Kentucky	34	-1	1029	3013	0.342	482	744	0.648	-1	...
1952	Kentucky	32	-1	1043	2829	0.369	549	865	0.635	-1	...
1953	Indiana	26	-1	737	2019	0.365	638	910	0.701	-1	...
1954	Kentucky	25	-1	829	2162	0.383	529	780	0.678	-1	...

	fg2_pct	fg3	fg3a	fg3_pct	opp_fg2	opp_fg2a	opp_fg2_pct	opp_fg3	\
1949	-1.0	-1	-1	-1.0	-1	-1	-1.0	-1	
1951	-1.0	-1	-1	-1.0	-1	-1	-1.0	-1	
1952	-1.0	-1	-1	-1.0	-1	-1	-1.0	-1	
1953	-1.0	-1	-1	-1.0	-1	-1	-1.0	-1	
1954	-1.0	-1	-1	-1.0	-1	-1	-1.0	-1	

	opp_fg3a	opp_fg3_pct
1949	-1	-1.0
1951	-1	-1.0
1952	-1	-1.0
1953	-1	-1.0
1954	-1	-1.0

[5 rows x 47 columns]

```
[8]: top_teams.tail()
```

	team_name	g	mp	fg	fga	fg_pct	ft	fta	ft_pct	orb	...	\
2016	Kansas	38	7700	1092	2207	0.495	601	843	0.713	402	...	
2017	Villanova	36	7200	965	1948	0.495	538	681	0.790	316	...	
2018	Virginia	34	6825	848	1844	0.460	340	451	0.754	282	...	
2019	Duke	38	7625	1157	2418	0.478	551	803	0.686	495	...	
2020	Kansas	31	6225	851	1758	0.484	411	616	0.667	333	...	

	fg2_pct	fg3	fg3a	fg3_pct	opp_fg2	opp_fg2a	opp_fg2_pct	opp_fg3	\
2016	0.533	304	728	0.418	640	1474	0.434	233	
2017	0.592	311	843	0.369	595	1211	0.491	251	
2018	0.501	247	645	0.383	432	1009	0.428	215	
2019	0.580	278	903	0.308	707	1571	0.450	253	
2020	0.553	199	578	0.344	455	1065	0.427	229	

	opp_fg3a	opp_fg3_pct
2016	726	0.321
2017	807	0.311
2018	694	0.310
2019	844	0.300
2020	750	0.305

[5 rows x 47 columns]

3.0.1 Data Processing Tutorial

Processing the data is another crucial step in our pipeline of analyzing this college basketball data. We attempted to use proper data science practice to make sure the `top_teams` pandas dataframe has labels and variable names that are accurate and concise and match the original dataset we scraped.

As you can see, all of the missing data has been encoded as -1, but the recent data is much more complete. In the original dataset, missing data was simply encoded with an empty string, so we used a simple regular expression to find the missing values. We have consciously decided to use -1 as the encoding for a missing value because it allows us to keep columns that should be integer values as such in our next step, not forcing them to be floating point numbers. This should save us some storage space. If we want to do any analysis on a column with missing data, we will just drop all of the columns that contain a -1 as that column's value before exploring it.

Finally, we convert all of the columns to their proper datatype. If a column involved a percentage or `per_game` value, then we encoded it as a float. Otherwise, the column's data was stored as a string. Excluding the `team_name` column, which is a string, all of our data is now in it's proper numeric datatype.

You can see that we've printed out the datatype of each column for convenience, as well as the head and tail of the dataset.

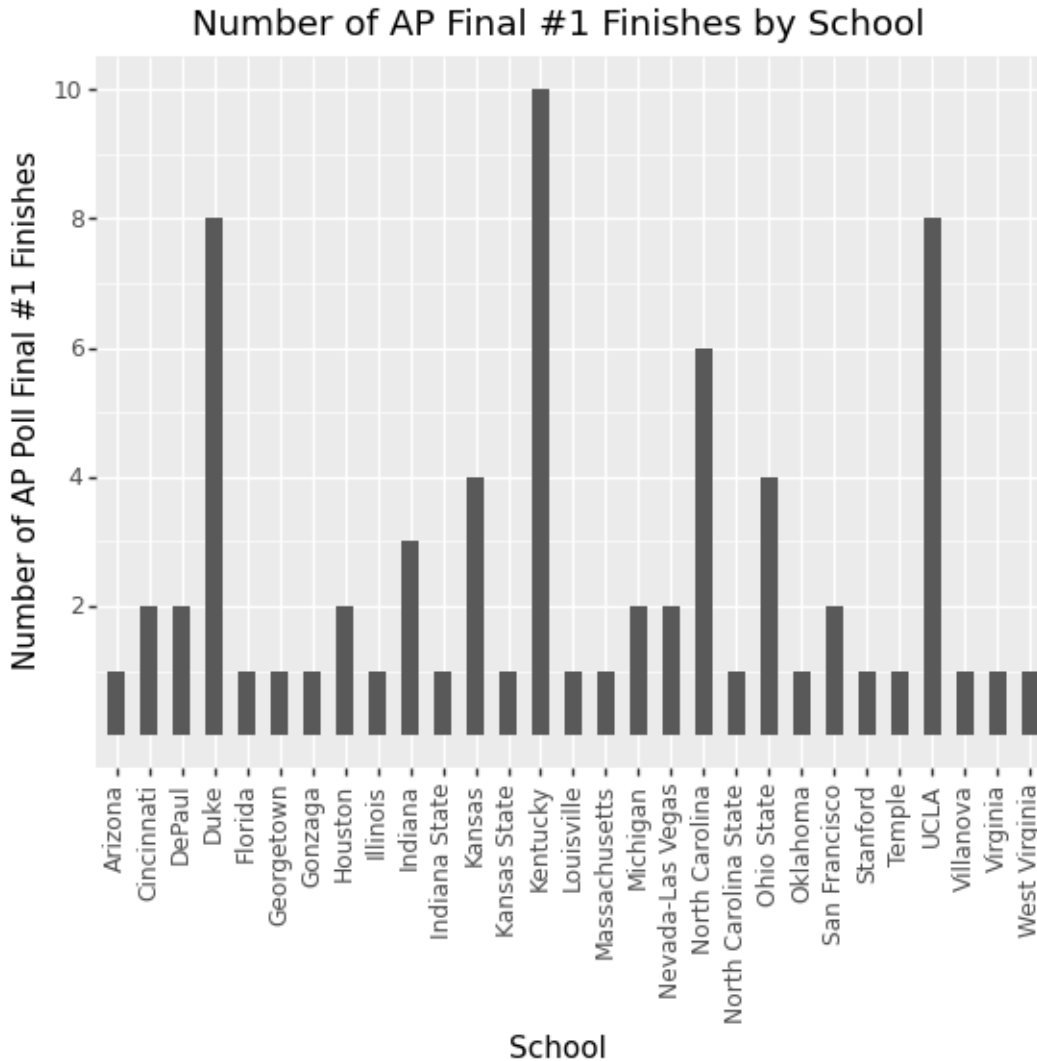
We are now ready to explore what makes a good NCAA basketball team.

4 Part 3 - Exploratory Data Analysis

Let's start by finding out which schools have been the most successful at producing the AP Poll #1 Ranked team at the end of each season.

```
[9]: # Let's see which colleges have been AP Poll #1 ranked the most often, after_
      ↪ each season
counts = top_teams.filter(['team_name', 'pts']).groupby(['team_name']).count()
counts['team_name'] = counts.index
counts.reset_index(drop = True, inplace = True)
counts.sort_values(by=['pts'], ascending = False, inplace = True)
```

```
[10]: (ggplot(counts, aes(x= 'team_name', y = 'pts')) +
      theme(axis_text_x = element_text(angle=90)) +
      xlab('School') +
      ylab('Number of AP Poll Final #1 Finishes') +
      ggtitle('Number of AP Final #1 Finishes by School') +
      scale_y_continuous(breaks=[2, 4, 6, 8, 10]) +
      geom_col(width = 0.5))
```



[10]: <ggplot: (300883741)>

Kentucky has achieved this feat a whopping 10 times. Duke and UCLA aren't far behind with 8 times, with North Carolina, Kansas, and Ohio State behind them at 4 times.


```
[11]: # Let's see how some key parts of basketball have changed over time. Let's
      ↪ start with points.
points = top_teams.filter(['team_name', 'pts_per_g', 'opp_pts_per_g'])
points['year'] = points.index.astype(int)
points.reset_index(drop = True, inplace = True)
points['differential'] = points['pts_per_g'] - points['opp_pts_per_g']
points.tail()
```



```
[11]:
```

	team_name	pts_per_g	opp_pts_per_g	year	differential
66	Kansas	81.3	67.6	2016	13.7
67	Villanova	77.2	62.7	2017	14.5
68	Virginia	67.1	54.0	2018	13.1
69	Duke	82.7	67.8	2019	14.9
70	Kansas	74.6	60.7	2020	13.9

```
[12]: (ggplot(points, aes(x= 'year', y = 'pts_per_g')) +
  geom_point() +
  theme(axis_text_x = element_text(angle=60)) +
  xlab('Year') +
  ylab('Points Per Game') +
  ggtitle('Points Per Game Scored over the Years') +
  geom_smooth(method = 'lm'))
```



```
[12]: <ggplot: (-9223372036551504240)>
```

```
[13]: import statsmodels.formula.api as sm
ppg_res = sm.ols('year~pts_per_g', data=points).fit()
ppg_res.summary()
```

```
[13]: <class 'statsmodels.iolib.summary.Summary'>
"""
```

```

                                OLS Regression Results
=====
Dep. Variable:                  year      R-squared:                0.020
Model:                            OLS      Adj. R-squared:            0.006
Method:                 Least Squares      F-statistic:                1.409
Date:                  Mon, 18 May 2020      Prob (F-statistic):          0.239
Time:                  13:52:13      Log-Likelihood:            -314.54
No. Observations:                71      AIC:                       633.1
Df Residuals:                    69      BIC:                       637.6
Df Model:                          1
Covariance Type:                nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2016.2214	26.426	76.295	0.000	1963.502	2068.941
pts_per_g	-0.3795	0.320	-1.187	0.239	-1.017	0.258

```

=====
Omnibus:                 13.220      Durbin-Watson:              0.030
Prob(Omnibus):            0.001      Jarque-Bera (JB):           3.928
Skew:                    -0.185      Prob(JB):                   0.140
Kurtosis:                 1.909      Cond. No.                   893.
=====

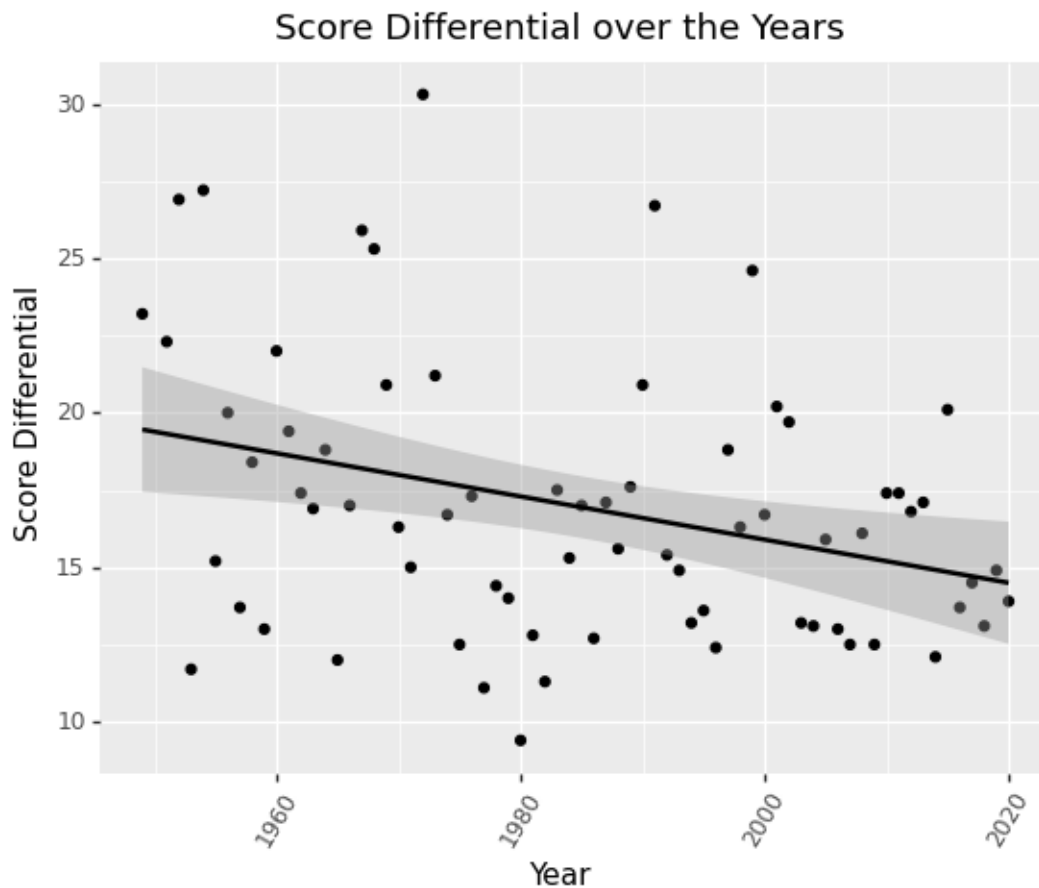
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
"""
```

From this plot, you can see that there may be a slight drop in points per game over the years by CBB's best teams. However, by checking the summary of this model, there is no statistically significant difference we can detect.

```
[14]: (ggplot(points, aes(x= 'year', y = 'differential')) +
geom_point() +
theme(axis_text_x = element_text(angle=60)) +
xlab('Year') +
ylab('Score Differential') +
ggtitle('Score Differential over the Years') +
geom_smooth(method = 'lm'))
```



```
[14]: <ggplot: (-9223372036551505306)>
```

```
[15]: diff_res = sm.ols('year~differential', data=points).fit()
diff_res.summary()
```

```
[15]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  year    R-squared:                0.106
Model:                            OLS    Adj. R-squared:           0.093
Method:                 Least Squares    F-statistic:                8.196
Date:                Mon, 18 May 2020    Prob (F-statistic):        0.00556
Time:                  13:52:13          Log-Likelihood:           -311.27
No. Observations:                  71    AIC:                       626.5
Df Residuals:                      69    BIC:                       631.1
Df Model:                           1
Covariance Type:                  nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2010.7460	9.296	216.297	0.000	1992.201	2029.292
differential	-1.5203	0.531	-2.863	0.006	-2.580	-0.461
=====						
Omnibus:		15.692	Durbin-Watson:			0.192
Prob(Omnibus):		0.000	Jarque-Bera (JB):			3.856
Skew:		-0.011	Prob(JB):			0.145
Kurtosis:		1.858	Cond. No.			69.9
=====						

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 ""

On the other hand, the best teams aren't winning by nearly as much anymore. This does have a statistically significant difference, providing a modeling equation of differential = 2010.7460 - 1.5203(year) with a p-value of .006. In context, this means that over the years, CBB games have gotten closer over time, and the best teams don't blow out their opponents nearly as often.

```
[16]: turnovers = top_teams.filter(['g', 'tov', 'opp_tov'])


# Some teams play more games, so we create a new statistic for both forced and
# given up turnovers per game played
turnovers['tov_per_g'] = turnovers['tov'] / turnovers['g']
turnovers['opp_tov_per_g'] = turnovers['opp_tov'] / turnovers['g']

# Let's also find the differential in turnovers.
turnovers['differential'] = turnovers['opp_tov_per_g'] - turnovers['tov_per_g']

# Convert year to a column
turnovers['year'] = turnovers.index.astype(int)
turnovers.reset_index(drop = True, inplace = True)

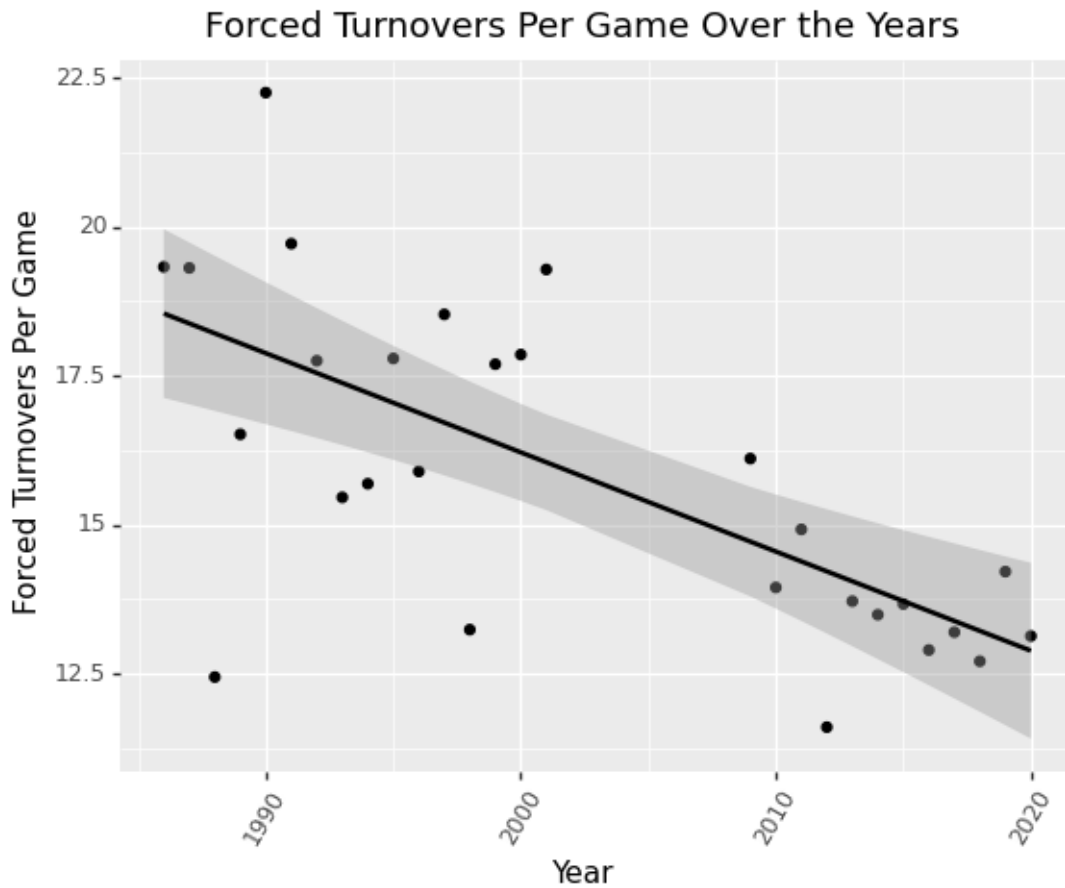
# Ignore all columns with missing values
turnovers = turnovers[(turnovers != -1).all(1)]

# Plot
(ggplot(turnovers, aes(x= 'year', y = 'tov_per_g')) +
geom_point() +
theme(axis_text_x = element_text(angle=60)) +
xlab('Year') +
ylab('Turnovers Per Game Given Up') +
ggtitle('Turnovers Per Game Given Up Over the Years') +
geom_smooth(method = 'lm'))
```



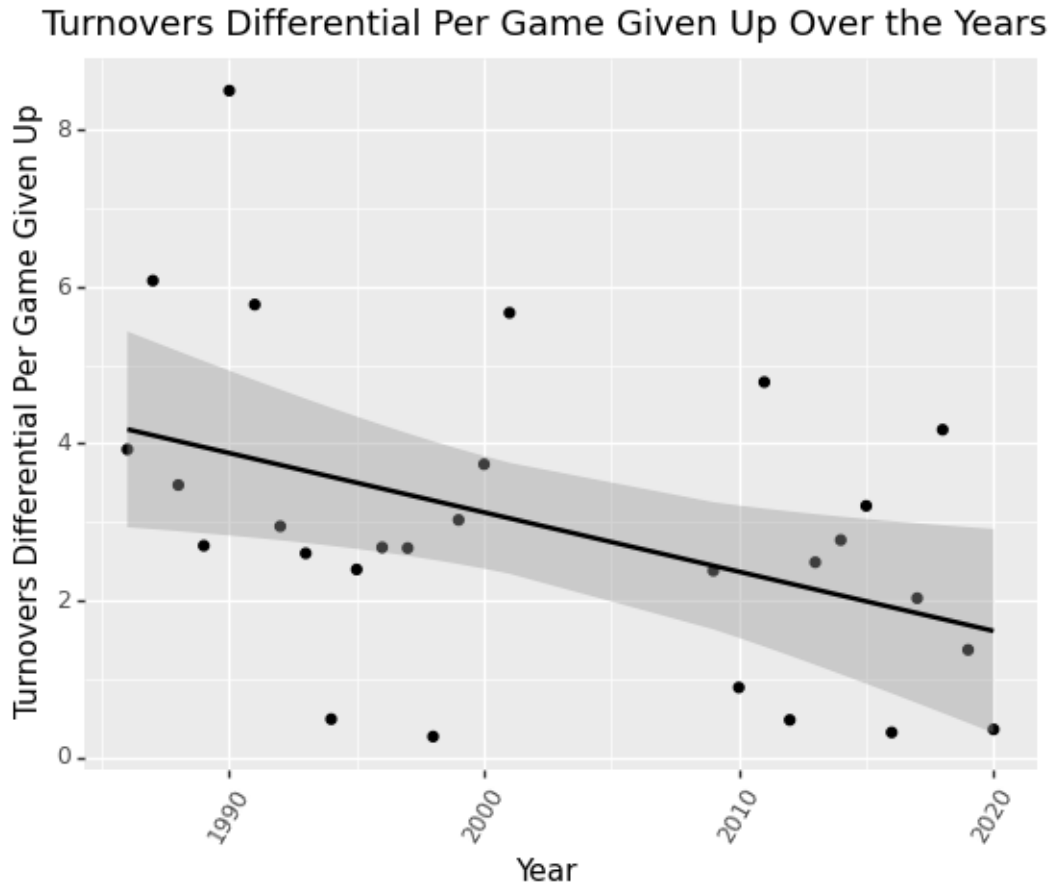
[16]: <ggplot: (303427497)>

```
[17]: (ggplot(turnovers, aes(x= 'year', y = 'opp_tov_per_g')) +
  geom_point() +
  theme(axis_text_x = element_text(angle=60)) +
  xlab('Year') +
  ylab('Forced Turnovers Per Game') +
  ggtitle('Forced Turnovers Per Game Over the Years') +
  geom_smooth(method = 'lm'))
```



[17]: <ggplot: (303278519)>

```
[18]: (ggplot(turnovers, aes(x= 'year', y = 'differential')) +
  geom_point() +
  theme(axis_text_x = element_text(angle=60)) +
  xlab('Year') +
  ylab('Turnovers Differential Per Game Given Up') +
  ggtitle('Turnovers Differential Per Game Given Up Over the Years') +
  geom_smooth(method = 'lm'))
```



```
[18]: <ggplot: (303644424)>
```

```
[19]: turnover_diff = sm.ols('year~differential', data=turnovers).fit()
      turnover_diff.summary()
```

```
[19]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  year    R-squared:                0.192
Model:                            OLS    Adj. R-squared:           0.161
Method:                 Least Squares    F-statistic:                6.166
Date:                Mon, 18 May 2020    Prob (F-statistic):          0.0198
Time:                  13:52:14    Log-Likelihood:            -104.38
No. Observations:                  28    AIC:                       212.8
Df Residuals:                      26    BIC:                       215.4
Df Model:                           1
Covariance Type:                  nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2009.9355	3.586	560.462	0.000	2002.564	2017.307
differential	-2.5350	1.021	-2.483	0.020	-4.634	-0.436
=====						
Omnibus:		7.974	Durbin-Watson:			0.354
Prob(Omnibus):		0.019	Jarque-Bera (JB):			2.120
Skew:		0.088	Prob(JB):			0.346
Kurtosis:		1.664	Cond. No.			6.75
=====						

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
"""
```

The change in turnover differential does have a statistically significant difference, providing a modeling equation of $\text{differential} = 2009.9355 - 2.5350(\text{year})$ with a p-value of .020. In context, this means that over the years, the best teams protect the ball on offense, knowing that they must not waste possessions by giving it away to the other team. The best teams outplay their opponents in the turnover categories, with the differential being > 0 , despite steadily dropping over the years as offenses protect the ball more.

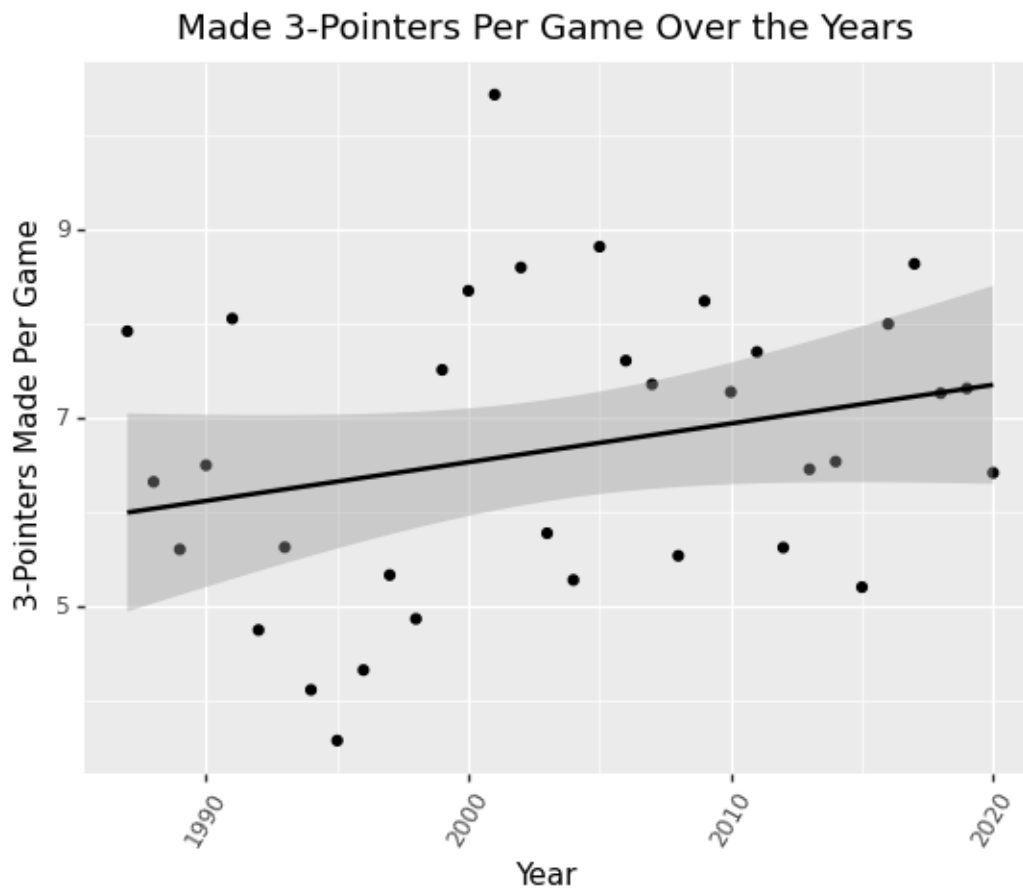
Turnovers have evolved into an important statistic in basketball. The best teams turn over the ball less often per game, but this also leads to their opponents turning the ball over less frequently too. We will see that this smaller number of possessions will lead to the rise of the 3-pointer (a way to score more points per possession than the standard 2 point shot) and promote a more efficient style of basketball.

```
[20]: threes = top_teams.filter(['g', 'fg3_pct', 'fg3', 'fg3a'])

# Some teams play more games, so we create a new statistic based on games played
threes['fg3_per_g'] = threes['fg3'] / threes['g']
threes['fg3a_per_g'] = threes['fg3a'] / threes['g']

# Convert year to a column
threes['year'] = threes.index.astype(int)
threes.reset_index(drop = True, inplace = True)
threes = threes[(threes != -1).all(1)]

# Plot
(ggplot(threes, aes(x= 'year', y = 'fg3_per_g')) +
 geom_point() +
 theme(axis_text_x = element_text(angle=60)) +
 xlab('Year') +
 ylab('3-Pointers Made Per Game') +
 ggtitle('Made 3-Pointers Per Game Over the Years') +
 geom_smooth(method = 'lm'))
```

```
[20]: <ggplot: (303752901)>
```

```
[21]: three_point_made = sm.ols('year~fg3_per_g', data=threes).fit()
three_point_made.summary()
```

```
[21]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  year    R-squared:                0.068
Model:                            OLS    Adj. R-squared:           0.039
Method:                 Least Squares    F-statistic:                2.332
Date:                Mon, 18 May 2020    Prob (F-statistic):          0.137
Time:                  13:52:15    Log-Likelihood:            -124.69
No. Observations:                  34    AIC:                       253.4
Df Residuals:                      32    BIC:                       256.4
Df Model:                            1
Covariance Type:                  nonrobust
=====
```


	coef	std err	t	P> t	[0.025	0.975]
-----	-----	-----	-----	-----	-----	-----
Intercept	1992.4627	7.419	268.548	0.000	1977.350	2007.576
fg3_per_g	1.6533	1.083	1.527	0.137	-0.552	3.859
=====	=====	=====	=====	=====	=====	=====
Omnibus:		2.848	Durbin-Watson:			0.093
Prob(Omnibus):		0.241	Jarque-Bera (JB):			1.390
Skew:		-0.033	Prob(JB):			0.499
Kurtosis:		2.012	Cond. No.			31.0
=====	=====	=====	=====	=====	=====	=====

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
"""
```

While the number of made three pointers per game has increased over the years since it's addition to the NCAA in 1986, the jump is not statistically significant with a p-value of 0.137.

```
[22]: (ggplot(threes, aes(x= 'year', y = 'fg3a_per_g')) +
geom_point() +
theme(axis_text_x = element_text(angle=60)) +
xlab('Year') +
ylab('3-Pointers Attempted Made Per Game') +
ggtitle('Attempted 3-Pointers Per Game Over the Years') +
geom_smooth(method = 'lm'))
```



```
[22]: <ggplot: (-9223372036551136985)>
```

```
[23]: three_point_att = sm.ols('year~fg3a_per_g', data=threes).fit()
three_point_att.summary()
```

```
[23]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  year    R-squared:                0.146
Model:                            OLS    Adj. R-squared:           0.120
Method:                 Least Squares    F-statistic:                5.490
Date:                Mon, 18 May 2020    Prob (F-statistic):        0.0255
Time:                  13:52:15          Log-Likelihood:           -123.19
No. Observations:                34      AIC:                   250.4
Df Residuals:                    32      BIC:                   253.4
Df Model:                        1
Covariance Type:                nonrobust
=====
```

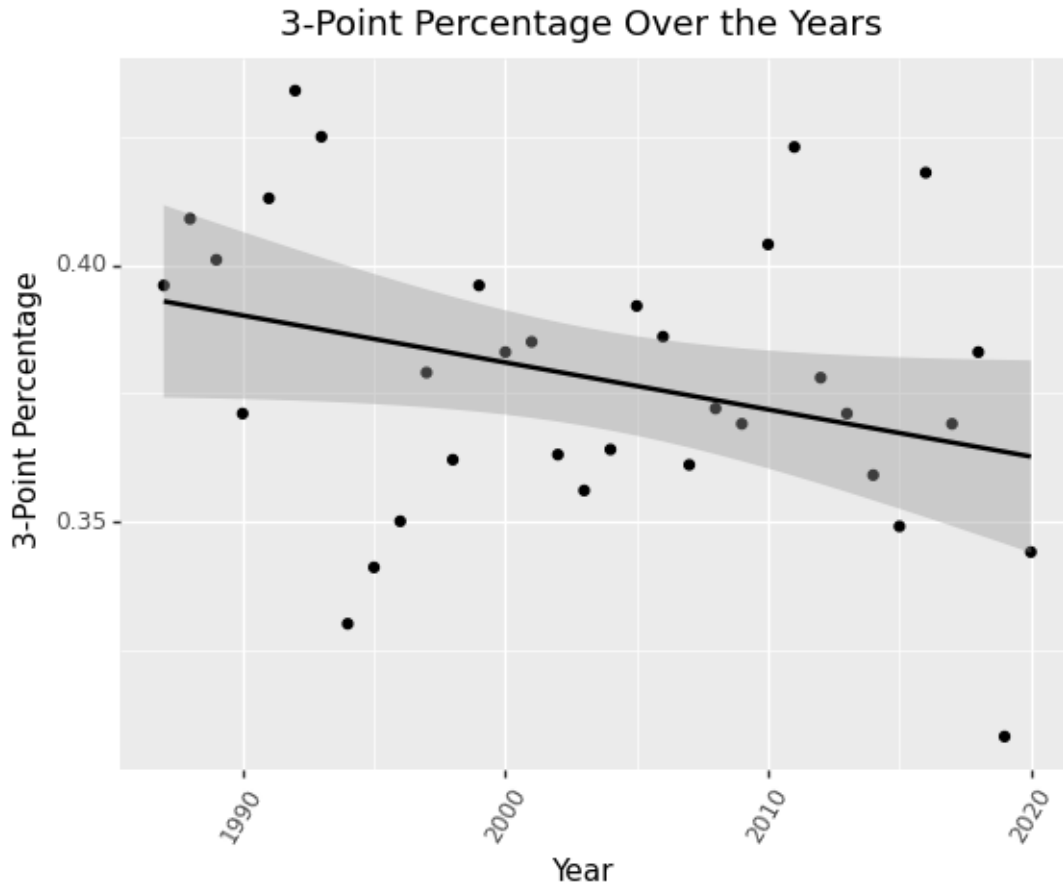
	coef	std err	t	P> t	[0.025	0.975]
Intercept	1986.8680	7.277	273.035	0.000	1972.045	2001.691
fg3a_per_g	0.9406	0.401	2.343	0.026	0.123	1.758
=====						
Omnibus:		2.603	Durbin-Watson:			0.196
Prob(Omnibus):		0.272	Jarque-Bera (JB):			1.362
Skew:		-0.097	Prob(JB):			0.506
Kurtosis:		2.039	Cond. No.			82.6
=====						

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
"""
```

On the other hand, the number of 3 point shots attempted has jumped in a statistically significant way, increasing by almost 1 for every year after 1987. The best teams in the NCAA have realized that the 3 pointer is the best ‘bang for your buck shot’ when compared to other jumpshots. Why shoot a jump shot from 20ft for 2 points, when you can step back another foot for a 3rd point? It only makes sense. The rising prominence of the three point shot changes the fundamental composition of basketball teams, making it less about height and bullying the other guy at the rim, and more about strategy and spacing of the court and good shooters.

```
[24]: (ggplot(threes, aes(x= 'year', y = 'fg3_pct')) +
  geom_point() +
  theme(axis_text_x = element_text(angle=60)) +
  xlab('Year') +
  ylab('3-Point Percentage') +
  ggtitle('3-Point Percentage Over the Years') +
  geom_smooth(method = 'lm'))
```



[24]: <ggplot: (303968022)>

It seems that the best teams are shooting more threes over time, yet making them at a worse percentage. Let's find the teams that were most efficient in their offenses. Since the possessions statistic is not available, we must calculate it using an estimation formula provided by <https://www.sportsrec.com/calculate-teams-offensive-defensive-efficiencies-7775395.html>

Possessions = Field Goals Attempted - Offensive Rebounds + Turnovers + (0.4 x free throws attempted)

```
[25]: # We will be calculating possessions per game
efficiency = top_teams.filter(['g', 'pts', 'opp_pts', 'fga', 'fta', 'orb',
    ↪ 'opp_drb', 'fg', 'tov', 'opp_fga', 'opp_FTA', 'opp_orb', 'drb', 'opp_fg',
    ↪ 'opp_tov', 'opp_fta', 'fg3a'])

# Every stat converted to its per game equivalency
efficiency = efficiency.div(efficiency['g'], axis=0)

# Add team name
```

```

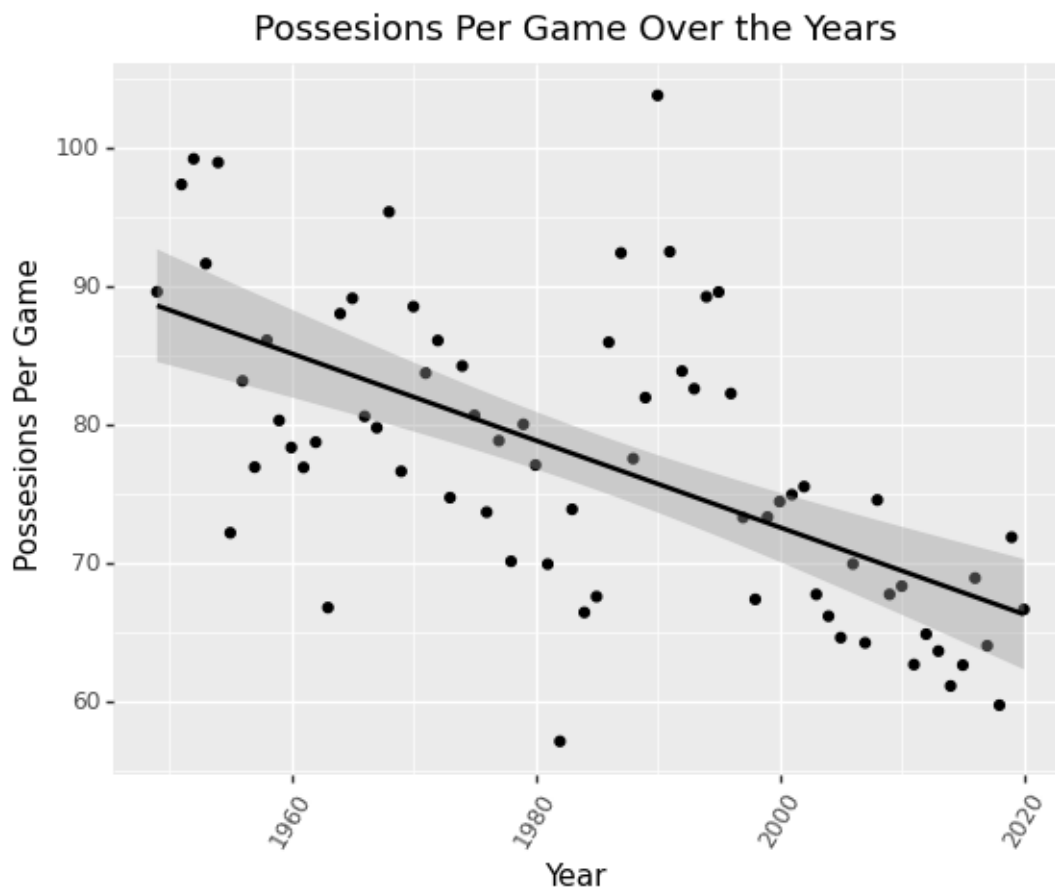
efficiency.insert(0, 'team_name', top_teams['team_name'])

# Estimating possessions
# field goals attempted - offensive rebounds + turnovers + (0.4 x free throws
  ↳ attempted) = total number of possessions
possessions = efficiency['fga'] - efficiency['orb'] + efficiency['tov'] + (0.4 *
  ↳ efficiency['fta'])

# Insert columns for possessions and year
efficiency.insert(1, 'possessions', possessions)
efficiency['year'] = efficiency.index.astype(int)

# Plot
(ggplot(efficiency, aes(x= 'year', y = 'possessions')) +
geom_point() +
theme(axis_text_x = element_text(angle=60)) +
xlab('Year') +
ylab('Possessions Per Game') +
ggtitle('Possessions Per Game Over the Years') +
geom_smooth(method = 'lm'))

```




[25]: <ggplot: (303171304)>

Now that we have a possessions estimate, we can calculate offensive efficiency for each team. The simple formula can be found on <https://www.nbastuffer.com/analytics101/offensive-efficiency/>

Offensive Efficiency Formula=(Points Scored)/(Possessions)

```
[26]: # Now we can calculate offensive efficiency
offensive = 100 * efficiency['pts'] / efficiency['possessions']
efficiency.insert(1, 'offensive', offensive)

# Plot
(ggplot(efficiency, aes(x= 'year', y = 'offensive')) +
geom_point() +
theme(axis_text_x = element_text(angle=60)) +
xlab('Year') +
ylab('Average Offensive Efficiency Rating') +
ggtitle('Average Offensive Efficiency Rating Over the Years') +
geom_smooth(method = 'lm'))
```



```
[26]: <ggplot: (-9223372036550925687)>
```

```
[27]: off_eff = sm.ols('year~offensive', data=efficiency).fit()  
off_eff.summary()
```

```
[27]: <class 'statsmodels.iolib.summary.Summary'>  
"""
```

```

                        OLS Regression Results
=====
Dep. Variable:          year      R-squared:                0.449
Model:                  OLS      Adj. R-squared:            0.441
Method:                 Least Squares      F-statistic:        56.28
Date:                  Mon, 18 May 2020    Prob (F-statistic):    1.62e-10
Time:                  13:52:16           Log-Likelihood:      -294.08
No. Observations:      71             AIC:                592.2
Df Residuals:          69             BIC:                596.7
Df Model:               1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	1849.9345	18.094	102.237	0.000	1813.837	1886.032
offensive	1.2557	0.167	7.502	0.000	0.922	1.590

```
=====
Omnibus:                8.157      Durbin-Watson:          0.339
Prob(Omnibus):          0.017      Jarque-Bera (JB):      2.983
Skew:                  0.113      Prob(JB):              0.225
Kurtosis:              2.022      Cond. No.              1.07e+03
=====
```

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.07e+03. This might indicate that there are strong multicollinearity or other numerical problems.

"""

```
[28]: efficiency[efficiency['offensive'] == efficiency['offensive'].max()]
```

```
[28]:   team_name  offensive  possessions    g      pts  opp_pts      fga \
1999      Duke  125.166026   73.358974  1.0  91.820513  67.153846  62.102564

      fta      orb  opp_drb      fg      tov  opp_fga \
1999  29.102564  15.051282  18.897436  31.897436  14.666667  63.74359
```


	opp_orb	drb	opp_fg	opp_tov	opp_fta	fg3a	year
1999	14.512821	27.128205	24.923077	17.692308	18.717949	18.948718	1999

As predicted, team strategies have evolved to be much more offensively efficient. The lack of extra possessions caused by an opponent's turnover has sped up this shift in play. The best teams typically gain 1.25 offensive efficiency rating points each year. The peak of offensive efficiency in our dataset was by the 1999 Duke team with an offensive efficiency rating of 125.17.


Now, let's do the same thing with defensive efficiency, calculating by the formula found on <https://www.nbastuffer.com/analytics101/defensive-efficiency/>

Defensive Efficiency = 100 * (Points Allowed / Possessions)

```
[29]: # Now let's calculate defensive efficiency
# Defensive Efficiency Formula=100*(Points Allowed/Possessions)

defensive = 100 * efficiency['opp_pts'] / efficiency['possessions']
efficiency.insert(1, 'defensive', defensive)

(ggplot(efficiency, aes(x= 'year', y = 'defensive')) +
 geom_point() +
 theme(axis_text_x = element_text(angle=60)) +
 xlab('Year') +
 ylab('Average Defensive Efficiency Rating') +
 ggtitle('Average Defensive Efficiency Rating Over the Years') +
 geom_smooth(method = 'lm'))
```



```
[29]: <ggplot: (-9223372036551264638)>
```

```
[30]: # The higher the efficiency the more points given up per 100 possesions, so a
      ↪ better team has a low rating
      efficiency[efficiency['defensive'] == efficiency['defensive'].min()]
```

```
[30]:      team_name  defensive  offensive  possessions    g      pts  opp_pts \
1949  Kentucky   50.210029  76.135469   89.623529   1.0  68.235294   45.0

           fga      fta      orb  ...      fg      tov  opp_fga \
1949  81.058824  21.411765 -0.029412 ...  26.558824 -0.029412  65.176471

           opp_orb      drb      opp_fg  opp_tov  opp_fta      fg3a  year
1949 -0.029412 -0.029412  15.823529    45.0  22.235294 -0.029412  1949

[1 rows x 21 columns]
```

```
[31]: def_eff = sm.ols('year~defensive', data=efficiency).fit()
      def_eff.summary()
```

```
[31]: <class 'statsmodels.iolib.summary.Summary'>
      """
                OLS Regression Results
=====
Dep. Variable:          year    R-squared:                0.459
Model:                  OLS    Adj. R-squared:           0.451
Method:                 Least Squares    F-statistic:            58.59
Date:                  Mon, 18 May 2020    Prob (F-statistic):      8.56e-11
Time:                  13:52:17    Log-Likelihood:         -293.44
No. Observations:      71    AIC:                    590.9
Df Residuals:          69    BIC:                    595.4
Df Model:               1
Covariance Type:       nonrobust
=====
                coef      std err          t      P>|t|      [0.025      0.975]
-----
Intercept    1873.5340     14.673    127.683     0.000    1844.262    1902.806
defensive      1.3028      0.170      7.654     0.000      0.963      1.642
=====
Omnibus:                 17.578    Durbin-Watson:           0.355
Prob(Omnibus):            0.000    Jarque-Bera (JB):        4.123
Skew:                    -0.089    Prob(JB):                0.127
Kurtosis:                 1.833    Cond. No.                 697.
=====

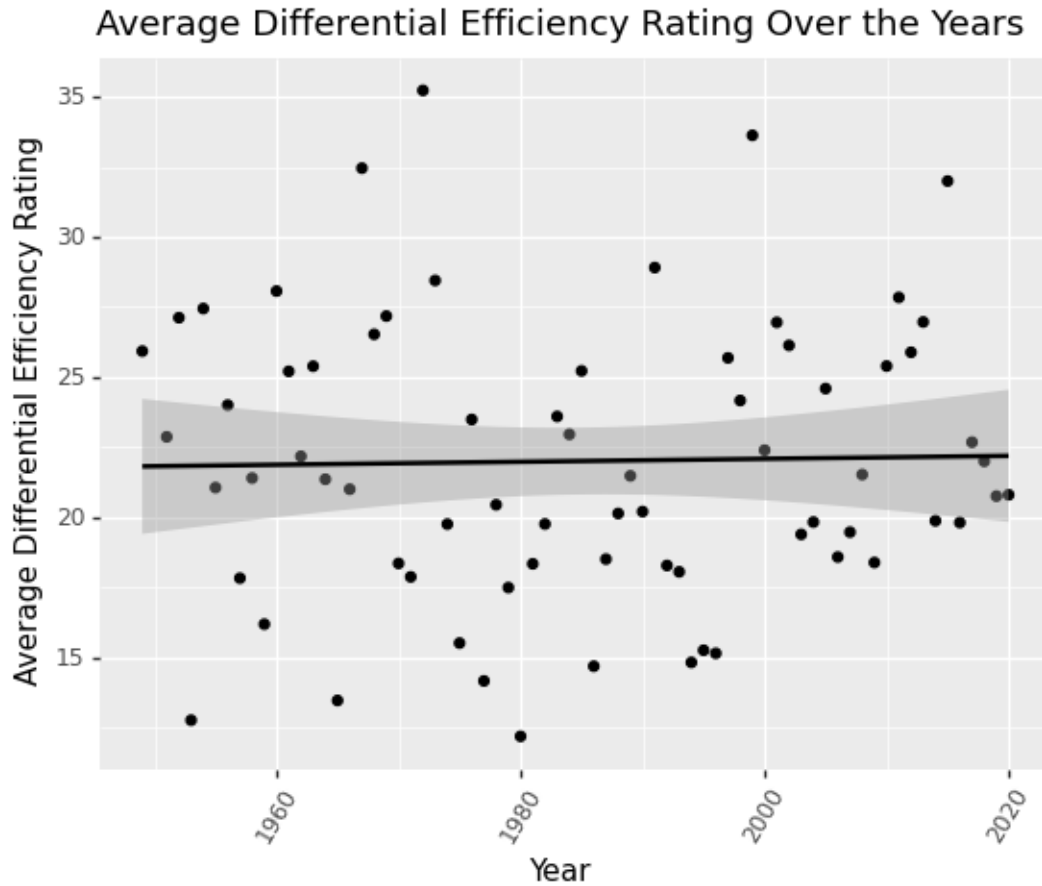
Warnings:
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
      """
```

Defensive efficiency is an interesting model because it has gotten significantly worse over the years, according to our simple regression model. Low numbers mean better defensive efficiencies here. We will attribute this shift to worse defensive efficiencies to the evolution of offensive tactics in today's game. Offenses are more efficient, so in response, defenses must be less efficient. Our previous plot about score differential showed that games are closer today than they were 15 years ago; it's not like defense doesn't matter at all anymore. Defenses simply aren't as dominant anymore as rules have been changed to put offenses in the game's spotlight.

```
[32]: # Since both offensive and defensive efficiencies are per points per 100
      ↪ possessions, a simple differential will tell us which teams perform on
      ↪ average on both sides of the court
differenital = efficiency['offensive'] - efficiency['defensive']
efficiency.insert(1, 'differential', differenital)

(ggplot(efficiency, aes(x= 'year', y = 'differential')) +
geom_point() +
theme(axis_text_x = element_text(angle=60)) +
```

```
xlab('Year') +
ylab('Average Differential Efficiency Rating') +
ggtitle('Average Differential Efficiency Rating Over the Years') +
geom_smooth(method = 'lm'))
```



```
[32]: <ggplot: (-9223372036551347596)>
```

```
[33]: # The most efficient team on both sides of the court
efficiency[efficiency['differential'] == efficiency['differential'].max()]
```

```
[33]:      team_name  differential  defensive  offensive  possessions    g  pts \
1972      UCLA      35.227625   74.63611  109.863735    86.106667  1.0  94.6

      opp_pts  fga      fta  ...  fg      tov  opp_fga  opp_orb \
1972  64.266667  75.4  26.766667  ...  38.0 -0.033333  66.766667 -0.033333

      drb      opp_fg  opp_tov  opp_fta      fg3a  year
1972 -0.033333  25.533333 -0.033333  19.266667 -0.033333  1972
```

```
[1 rows x 22 columns]
```

Interestingly, this team is widely considered one of the greatest of all time.
<https://bleacherreport.com/articles/1046550-the-50-best-teams-in-college-basketball-history>

```
[34]: # The least efficient team on both sides of the court
      efficiency[efficiency['differential'] == efficiency['differential'].min()]
```

```
[34]:      team_name  differential  defensive  offensive  possessions  g  \
1980    DePaul      12.180437  96.470915  108.651352    77.114286  1.0

      pts  opp_pts  fga      fta  ...      fg      tov  \
1980  83.785714  74.392857  67.5  24.035714  ...  32.821429 -0.035714

      opp_fga  opp_orb      drb      opp_fg  opp_tov  opp_fta      fg3a  \
1980  70.178571 -0.035714 -0.035714  31.392857 -0.035714  16.285714 -0.035714

      year
1980  1980
```

```
[1 rows x 22 columns]
```

While the efficiency differential doesn't really show a trend, it does separate the best of the best teams. The 1972 UCLA Bruins had the best efficiency differential in our dataset, and went on to win the 1972 National Championship. More recently, the 2015 Kentucky team that was extremely efficient with a differential of ~32 made a Final Four run. On the other hand, the 1980 DePaul team had an awful efficiency differential, and was eliminated from the NCAA tournament in the Round of 32.

4.0.1 Exploratory Data Analysis Tutorial

After collecting and processing the dataset, the next step is to visualize it. We must be mindful of what we are plotting to avoid the missing data values.

Out of pure curiosity, we checked to see which schools produced the best teams most often, finding that Kentucky has produced 10 AP Poll #1 teams. Duke and UCLA aren't far behind with 8 teams apiece, and North Carolina, Kansas, and Ohio State behind them at 4 teams in our dataset.

Next, we filtered out rows with -1 values, and plotted trends in points, score differential, turnovers (and differential), 3 point shooting, and offensive / defensive efficiencies. In addition to scatter plots, we overlayed a linear regression line and printed the linear model for each of our inquiries.

Finally, we added a bit of prose between each inquiry to see what we could draw from our data analysis of those specific basketball statistic categories.

Some statistics didn't show much, but others such as score differential, attempted three point shooting, and efficiency showed statistically significant trends that inform us on how NCAA basketball has evolved over the last 60+ years. From these trends, we can infer what is important to today's NCAA game, such as the run and gun 3 point shooting offenses, and efficient defenses.

5 Part 4 - Hypothesis Testing and Machine Learning

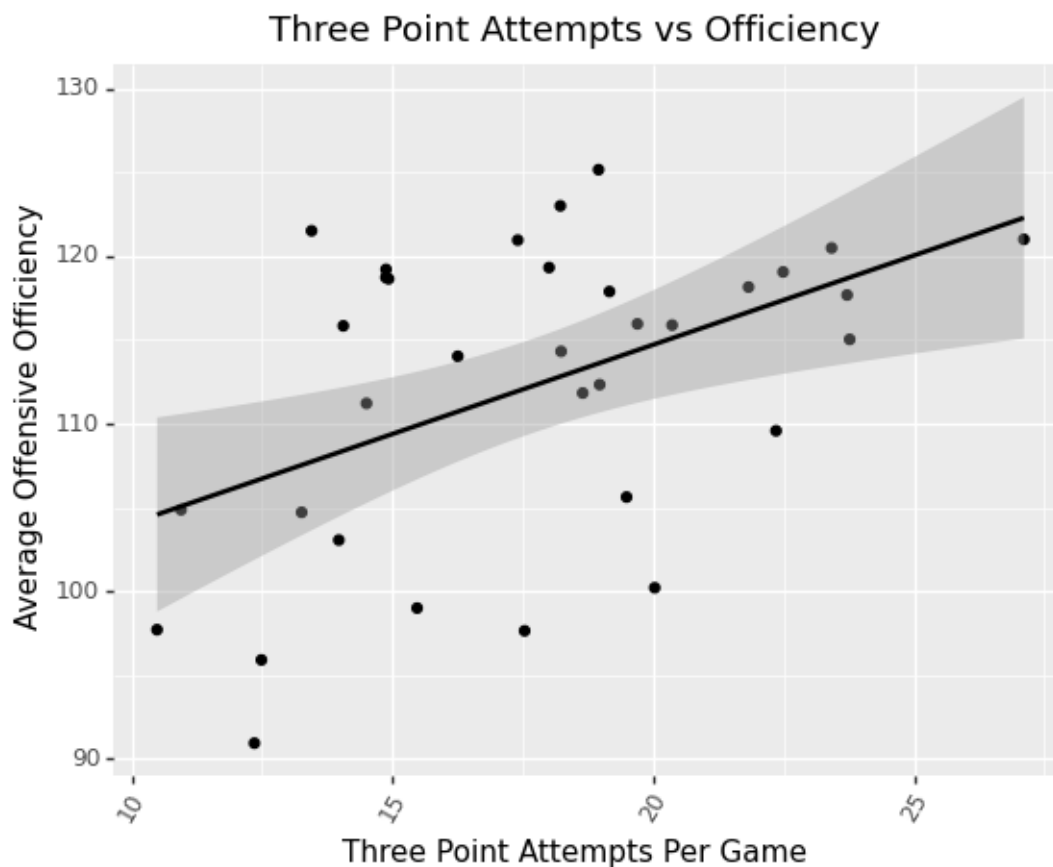
If you haven't noticed, every scatter plot we've produced is accompanied by line of linear regression. It helps to establish certain trend indicators across time, and it is also the simplest form of machine learning. See, if you look at the offensive efficiency over time graph, you could extrapolate that in the future, teams will be more efficient in their scoring. Linear regression is simply a technique to minimize the cost function and the residuals, between our model and the data.

Let's create a new problem with our own hypothesis, one not involving time. It is often speculated in the new age of basketball, shooting more three point shots makes an offense more efficient. You can read more about this so-called "revolution" in basketball, and specifically college ball, here <https://bleacherreport.com/articles/2762158-the-3-point-revolution-has-taken-over-college-basketball-too>.

Our null hypothesis will be that there is no association between how a team's offensive efficiency and the number of threes they attempt. The alternate hypothesis is that there is a correlation between offensive efficiency and the number of attempted threes.

```
[35]: # First lets remove the columns with missing values, since the three point line
      ↪was introduced to NCAA in 1986
      efficiency.drop(efficiency[efficiency['fg3a'] < 0].index, inplace = True)

      # Plot regression
      (ggplot(efficiency, aes(x= 'fg3a', y = 'offensive')) +
       geom_point() +
       theme(axis_text_x = element_text(angle=60)) +
       xlab('Three Point Attempts Per Game') +
       ylab('Average Offensive Efficiency') +
       ggtitle('Three Point Attempts vs Efficiency') +
       geom_smooth(method = 'lm'))
```



[35]: <ggplot: (303750720)>

```
[36]: model = sm.ols('fg3a~offensive', data=efficiency).fit()
      model.summary()
```

[36]: <class 'statsmodels.iolib.summary.Summary'>
 """

```

                                OLS Regression Results
=====
Dep. Variable:                  fg3a    R-squared:                0.229
Model:                            OLS    Adj. R-squared:            0.205
Method:                 Least Squares    F-statistic:                9.522
Date:                Mon, 18 May 2020    Prob (F-statistic):          0.00417
Time:                  13:52:18    Log-Likelihood:             -90.877
No. Observations:                  34    AIC:                        185.8
Df Residuals:                      32    BIC:                        188.8
Df Model:                           1
Covariance Type:                nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-6.4622	7.849	-0.823	0.416	-22.450	9.526
offensive	0.2151	0.070	3.086	0.004	0.073	0.357
=====						
Omnibus:		2.431	Durbin-Watson:			1.432
Prob(Omnibus):		0.297	Jarque-Bera (JB):			1.465
Skew:		0.223	Prob(JB):			0.481
Kurtosis:		2.086	Cond. No.			1.43e+03
=====						

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.43e+03. This might indicate that there are strong multicollinearity or other numerical problems.

"""

With this model we can see a linear regression model that fairly closely matches are expected result, that three-point attempts do correlate with offensive efficiency. If we look closely at our model, which utilizes the least squares method, has an R-squared value of 0.229 which is pretty good. The p-value which is used to accept or reject our null hypothesis, is very low at 0.004, much less than our needed alpha level of significance of 0.05. This means, we can reject our null hypothesis of no correlation between offensive efficiency and three-point attempts. There is strong evidence that such a correlation between the two statistics exists.

Then by extrapolating further, this linear regression model can further be used as predictive measure to show the predicted offensive efficiency based on the number of three's attempted by top ranked college basketball teams.

6 Part 5 - Final Thoughts

We hope you found this tutorial interesting and informative. Data science and sports go hand in hand, nowadays. It has become increasingly clear that professional, collegiate, and amateur sports teams can all utilize increased data analytics to improve their respective games. As found with this tutorial, top collegiate teams improve each coming year, and so do their opponents.

Gaining a competitive edge is not an easy task. Teams can you use machine learning to best predict the most effective strategies, in our model, that was shooting more three pointers. If you want to do more research in this field, much has been written about data analytics in the NBA, and it can all be generalized to all levels of basketball.

Here is some more information:

<https://www.nytimes.com/2019/11/27/sports/basketball/nba-analytics.html>

These ones is all about the three-point shot:

<https://towardsdatascience.com/nba-data-analytics-changing-the-game-a9ad59d1f116>

<https://onlinedsa.merrimack.edu/nba-analytics-changing-basketball/>

[]: